

02 - N-ATLaS integration evidence

PauseAm - Ask before you pay. Innovation & Enterprise / PS2. Review copy, 3 October 2026. Intended owner review/submission: Monday, 5 October. Not submitted.

Latest demonstrated release

Version 25 is live with English ASR, pinned N-ATLaS CPU relevance inference and device-read spoken replies with text below and replay/stop. Typed replies stay silent. All 80 application tests and ten final public backend checks passed; four backend checks used the genuine N-ATLaS text model. The owner confirms the full voice journey and controls work. These are engineering/owner checks, not participant validation. Current live source commit: 37a88225905963275b65eedb1ebf28bdf286f065. Deployment: appgdep_6ac0abcd1a148191a1f4be4f52a4def1; environment revision 4.

Latest release evidence: <https://github.com/thepopeblack-byte/pauseam/blob/main/submission/evidence/spoken-replies-release-2026-10-03.json> Latest genuine request traces: <https://github.com/thepopeblack-byte/pauseam/blob/main/submission/evidence/spoken-replies-live-2026-10-03.json>

The owner reports testing is under way, but no completed observer records have been supplied. Documented participant count remains 0. Official-service qualification remains unresolved. Other languages, official API and fine-tuning are not claimed. Older request outcomes below retain their historical dates.

Current capability boundary

Amount/date and bank-help addendum: corrected/typed payment questions accept ordinary amounts and dates; number values are omitted before the genuine model relevance request. The active ASR continues to hide numbers. An independent official-source directory answers email, phone and basic USSD-menu questions for eight Nigerian banks without model inference (model=null). These contacts are published bank facts, not model generation or official N-ATLaS API integration. See [../evidence/bank-information-source-review-2026-10-03.md](#). Independent human source review and documented participant results remain pending.

Historical 3 October checkpoint: Sites version 20 was live. The existing VM's authenticated ASR health reports the pinned English model and numbers-redacted-v1 privacy policy; the pinned llama.cpp text service is ready. Public authored questions with a redacted amount returned exact reviewed supplier and reporting cards. A silent WAV received the distinct 422 recording error without a transcript. Actual timestamps, identities and latency are in [../evidence/freeform-voice-live-2026-10-03.json](#). The owner confirms that the unscripted amount/date full voice journey works. This is a self-reported check, not independent observation, speech accuracy evidence or a completed participant interaction.

The agreed public language is Nigerian-accented English. The official English ASR returned a real owner transcript through the app; this is engineering evidence, not a completed participant-validation round. After the owner restored missing encrypted configuration, the correct same-model CPU service became ready. Six genuine authenticated HTTPS relevance requests passed in 6,141-6,754 ms, plus two pre-inference guards. Sites version 19 was redeployed with TEXT_ENABLED=true at 22:36:47 UTC on 2 October. Public model-backed supplier guidance was verified in the actual browser. These are engineering checks, not human validation.

Yoruba, Hausa and Igbo adapters and byte-verified model files remain available but are paused. Official API access and fine-tuning are not demonstrated. Quantization is a precision conversion, not fine-tuning. Do not select unsupported capabilities or languages on the portal.

Actual architecture

Consent -> English recording -> bounded mono 16 kHz WAV -> Sites Worker -> authenticated SecretVM ASR -> identity/privacy checks -> transcript correction/confirmation -> current source cards -> optional N-ATLaS card selection -> independently validated checklist.

Urgent bank-first actions do not wait for a model. Reports are reviewed drafts, not automatically sent complaints. No funds are moved, accounts authenticated, safe verdicts issued or recovery promised.

Public app: <https://pauseam.theblockcapitol.com>

Public source: <https://github.com/thepopeblack-byte/pauseam>

Team inference: <https://amaranth-nightingale.vm.scrtlabs.com>

Model identities and recipe

English ASR: NCAIR1/NigerianAccentedEnglish at 3c52c6e6c9ec508014a7b9db6a42b503b8930dff. Its separate approved bucket is Blockcapitol/NigerianAccentedEnglish-bucket; every required file was checked against that official revision. See `asr/manifests/en.json` and `submission/evidence/asr-bucket-en-2026-10-02.json`.

Text: NCAIR1/N-ATLaS at e294476928aca9030e924ca27bb8e085e8581273. The new CPU path verifies original Safetensors/tokenizer bytes against `text_service/model-manifest.json`, converts with official `llama.cpp` revision 631109b34da437a3c4a5ebd75091d677671392e3 and quantizes to Q4_K_M. The embedded official chat template is used. The generated provenance receipt records original hashes, output hash/size and recipe. No alternative model is substituted.

Runtime release b11351 is pinned. Linux archive SHA256:

c800a3402548d57f408adcebe4e53b939141e34e09b7e4701c785435572d1fee. Downloading files alone proves neither inference nor quality.

Code locations

`lib/asr.ts`, `app/api/asr/route.ts` and `asr/app.py` enforce consent, bounded audio, exact official identity and sensitive-output refusal. `lib/text-model.ts` validates the selected card and optional CPU runtime identity. `app/api/answer/route.ts` keeps urgent source content immediate and fails visibly if enabled inference/configuration is unavailable.

`text_cpu/bootstrap.py` verifies original files and the conversion cache. `text_cpu/contract.py` creates finite schema-constrained generation; `text_cpu/app.py` authenticates, serializes and reserves real inference. Source retrieval chooses the strongest current card; N-ATLaS checks its relevance and may accept or abstain. Only that card ID or an empty selection is accepted. Unrelated/instruction-attack inputs are stopped before inference. Model-authored contacts, links and arbitrary safety prose never render. This is constrained relevance checking, not open-ended conversational generation.

`scripts/prepare-secretvm.py` generates digest-pinned profiles. `scripts/verify-model-host.mjs` checks real identities/requests. The old full-precision `text_service` and prior traces remain as historical evidence.

Redacted production settings

ASR_ENABLED=true; ASR_ENDPOINT=<https://amaranth-nightingale.vm.scrtlabs.com/asr/en/transcribe>;
KB_ENABLED=true; TEXT_ENABLED=true;
TEXT_ENDPOINT=<https://amaranth-nightingale.vm.scrtlabs.com/text/guide>. ASR_SERVICE_TOKEN and TEXT_SERVICE_TOKEN are server-side secrets, excluded here. HF_TOKEN is private to model hosting/preparation. Deployment `appgdep_6ac031f10d708191a2bb327e3af7dc14` applied environment revision 4.

Text is enabled only after genuine host identity, relevance, adversarial and latency checks, followed by a Sites deployment. Missing settings or failed inference cannot be passed off as model guidance.

Genuine evidence and gaps

The earlier English owner transcript is documented in `worker-redirect-fix-2026-10-02.json`. The latest authenticated English health check returned HTTP 200 in 835 ms with the exact identity. Provenance/privacy/relevance changes passed 56 application tests and 27 Python checks. These are not participant interactions or speech-accuracy measurements.

Local conversion and successful requests remain in `cpu-conversion-2026-10-02.json` and `cpu-local-relevance-check-2026-10-02.json`. Earlier all-card/three-candidate failures remain unchanged. A schema-wrapper defect was caught and fixed before enablement. Retrieval followed by relevance checking passed; model-only classification of all cards did not.

Historical failures remain unchanged. Final container proof is `cpu-final-container-2026-10-02.json`; actual build: <https://github.com/thepopeblack-byte/pauseam/actions/runs/37061753954>. Its digest starts `ed76b851`; live readiness identifies `llama.cpp` and `reviewed-card-relevance-v1`. The VM-generated quantized SHA256 is `3820854be929790f10d171cd6f20dcd4e1ab3ccba095c133144a8dd10d65e49b`; the separately generated Windows conversion has a different hash. Original source files are pinned in both recipes. Actual request trace: <https://github.com/thepopeblack-byte/pauseam/blob/main/submission/evidence/live-text-inference-2026-10-02.json>

Historical public results: `live-text-public-check-2026-10-02.json`. Of 21 public checks, 20 passed. The vague question "They want me to use different bank details" returned a safe no-match; the explicit supplier question worked. This historical outcome is retained. Version 24 subsequently added clarification for vague payment concerns; representative comprehension testing remains pending. Five additional reporting/learning checks passed. The owner reports that the fresh complete English recording, transcript review and supplier guidance journey works; this is owner confirmation, not independently observed participant validation. Random response references are not signed proof of human participation: observed sessions require attestation.

Official service and API findings

Version 24 adds source-derived situation explanations and clarifying questions; the model's constrained card-selection contract and original source card remain unchanged. It does not claim that authored explanations were generated by N-ATLaS. Four genuine final public requests carried the pinned text-model identity and returned eligible cards; see `evidence/contextual-replies-live-v24-2026-10-03.json`. Vague payment queries request clarification without a payment assessment or claimed model inference. Version 25 adds automatic device readout after a corrected, confirmed voice question, with reply text below and replay/stop controls. Typed questions remain silent. N-ATLaS ASR and text inference are retained; no official N-ATLaS TTS model/API was found in its published model inventory and licence. Readout uses local browser/device speech, not another advice model or an OpenAI API. The owner confirms voice, replay/stop and silent text work. This is one owner check, not a recorded research interaction. See `../evidence/spoken-replies-implementation-2026-10-03.md` and `../evidence/spoken-replies-release-2026-10-03.json`. This review copy includes the spoken-reply update; final human validation and submission review remain pending.

PS2 requires the official N-ATLaS ASR service for the relevant language: <https://ncair.nitda.gov.ng/naic/>

Organiser acceptance of self-hosted official weights remains pending. Kayode authorized the prepared secretariat query, but no successful dispatch or reply is evidenced here. The official page lists API credentials among shortlisted-team support; no verified official endpoint/contract or successful official request is established. `lib/official-api.ts` fails closed. Official API and fine-tuning checkboxes remain unverified.

Licence

The published cap is 1,000 active end users within a rolling 30 days; obtain separate licensing before exceeding it. The published scope also requires a separate agreement for commercial use, so obtain paid-use permission before commercial launch even below that cap. Model access approval is not proof of commercial permission. This deployment uses a shared persistent ceiling of 950 inference reservations, including failures. Preserve its ledger across updates and keep alternate hosts inactive. Model terms: <https://huggingface.co/NCAIR1/N-ATLaS>

N-ATLaS is an initiative of the Federal Ministry of Communications, Innovation and Digital Economy, and powered by Awarri Technologies.

Actual public text journey

2 October 2026. An authored supplier question returned the reviewed checklist through the public app after genuine N-ATLaS inference was enabled. This screenshot is operating-product evidence, not a participant result. The owner separately confirms the complete English voice journey works.

 **PauseAm**
Ask before you pay.

How it works

Ask a question

Report a problem

Learn

What's happening?

Ask about a payment, message or money concern. Use your own words.

A supplier sent new bank details and wants me to pay today.

 Speak instead

Get next steps

 Leave out names, numbers, PINs, OTPs and passwords.

✓ Your next steps

Verify changed supplier details independently

- 1** Pause this payment.
A changed instruction needs a fresh check.
- 2** Call your supplier using a contact established before this request.
Confirm the change separately.
A reply to the same message may reach the impersonator.
- 3** If you cannot confirm the request independently, keep the payment on hold.
Uncertainty is not evidence that a payment is safe.

Share checklist

 Listen

[Ask another question](#)

English · No account needed

Privacy

Current mobile interface - version 25

3 October 2026. Actual operating product in a 390-pixel emulated viewport, showing the Speak / Type interface. This screenshot proves the interface state; it does not prove speech accuracy, physical-device performance or participant validation. The separate genuine traces and owner check are linked in this report.

