

02 - N-ATLaS integration evidence

PauseAm - Ask before you pay. Innovation & Enterprise / PS2. Review copy, 2 October 2026. Intended owner review/submission: Monday, 5 October. Not submitted.

Current capability boundary

The agreed public language is Nigerian-accented English. The official English ASR returned a real owner transcript through the app; this is engineering evidence, not a completed participant-validation round. The new same-model CPU service completed six genuine local relevance requests in 15,909-23,293 ms, plus two source/input guard checks without inference. Its actual conversion output is 4,920,738,848 bytes. Public text inference remains disabled until the owner updates SecretVM and actual HTTPS requests pass. Local success is not VM success.

Yoruba, Hausa and Igbo adapters and byte-verified model files remain available but are paused. Official API access and fine-tuning are not demonstrated. Quantization is a precision conversion, not fine-tuning. Do not select unsupported capabilities or languages on the portal.

Actual architecture

Consent -> English recording -> bounded mono 16 kHz WAV -> Sites Worker -> authenticated SecretVM ASR -> identity/privacy checks -> transcript correction/confirmation -> current source cards -> optional N-ATLaS card selection -> independently validated checklist.

Urgent bank-first actions do not wait for a model. Reports are reviewed drafts, not automatically sent complaints. No funds are moved, accounts authenticated, safe verdicts issued or recovery promised.

Public app: <https://pauseam.theblockcapitol.com>

Public source: <https://github.com/thepopeblack-byte/pauseam>

Team inference: <https://amaranth-nightingale.vm.scrtlabs.com>

Model identities and recipe

English ASR: NCAIR1/NigerianAccentedEnglish at 3c52c6e6c9ec508014a7b9db6a42b503b8930dff. Its separate approved bucket is Blockcapitol/NigerianAccentedEnglish-bucket; every required file was checked against that official revision. See `asr/manifests/en.json` and `submission/evidence/asr-bucket-en-2026-10-02.json`.

Text: NCAIR1/N-ATLaS at e294476928aca9030e924ca27bb8e085e8581273. The new CPU path verifies original Safetensors/tokenizer bytes against `text_service/model-manifest.json`, converts with official `llama.cpp` revision 631109b34da437a3c4a5ebd75091d677671392e3 and quantizes to Q4_K_M. The embedded official chat template is used. The generated provenance receipt records original hashes, output hash/size and recipe. No alternative model is substituted.

Runtime release b11351 is pinned. Linux archive SHA256:

c800a3402548d57f408adcebe4e53b939141e34e09b7e4701c785435572d1fee. Windows archive SHA256:

ced25d91ed2c0981420dcfc9e7823892156c13b5c8dc9f26d334d8e3e4c33a09. Conversion/quality/latency are not assumed from downloading these files.

Code locations

`lib/asr.ts`, `app/api/asr/route.ts` and `asr/app.py` enforce consent, bounded audio, exact official identity and sensitive-output refusal. `lib/text-model.ts` validates the selected card and optional CPU runtime identity. `app/api/answer/route.ts` keeps urgent source content immediate and fails visibly if enabled inference/configuration is unavailable.

`text_cpu/bootstrap.py` verifies original files and the conversion cache. `text_cpu/contract.py` creates finite schema-constrained generation; `text_cpu/app.py` authenticates, serializes and reserves real inference. Source retrieval chooses the strongest current card; N-ATLaS checks its relevance and may accept or abstain. Only that card

ID or an empty selection is accepted. Unrelated/instruction-attack inputs are stopped before inference. Model-authored contacts, links and arbitrary safety prose never render. This is constrained relevance checking, not open-ended conversational generation.

scripts/prepare-secretvm.py generates digest-pinned profiles. scripts/verify-model-host.mjs checks real identities/requests. The old full-precision text_service and prior traces remain as historical evidence.

Redacted production settings

ASR_ENABLED=true; ASR_ENDPOINT=https://amaranth-nightingale.vm.scrtlabs.com/asr/en/transcribe;
KB_ENABLED=true; TEXT_ENABLED=false at this checkpoint;
TEXT_ENDPOINT=https://amaranth-nightingale.vm.scrtlabs.com/text/guide. ASR_SERVICE_TOKEN and
TEXT_SERVICE_TOKEN are server-side secrets, excluded here. HF_TOKEN is private to model hosting/preparation.

Text is enabled only after genuine host identity, relevance, adversarial and latency checks, followed by a Sites deployment. Missing settings or failed inference cannot be passed off as model guidance.

Genuine evidence and gaps

The English owner transcript and redirect investigation are documented in worker-redirect-fix-2026-10-02.json and english-pilot-2026-10-02.json. This turn's English authenticated health check returned HTTP 200 in 809 ms with the exact identity; no new audio was uploaded. New provenance/privacy/relevance changes passed 56 application tests and 27 Python checks. These are engineering checks, not participant interactions or ASR accuracy measurements.

Actual local conversion: submission/evidence/cpu-conversion-2026-10-02.json. Actual final local outcomes: submission/evidence/cpu-local-relevance-check-2026-10-02.json. Earlier all-card and three-candidate failures remain in cpu-local-headings-first-check-2026-10-02.json and cpu-three-candidate-check-2026-10-02.json. A required runtime schema-wrapper bug was caught by output validation and fixed before enabling public inference. The final pipeline uses retrieval followed by a model relevance check; it does not assert that model-only classification of all cards passed.

Historical failures remain unchanged. Final container proof is cpu-final-container-2026-10-02.json; actual build: <https://github.com/thepopeblack-byte/pauseam/actions/runs/37061753954>. Its digest starts ed76b851 and readiness must identify reviewed-card-relevance-v1. A build/import check is not inference. Random response references are not signed proof of human participation: observed sessions require attestation. Final live requests and screenshots remain pending.

Official service and API findings

PS2 requires the official N-ATLaS ASR service for the relevant language: <https://ncair.nitda.gov.ng/naic/>

Organiser acceptance of self-hosted official weights remains pending. Kayode authorized the prepared secretariat query, but no successful dispatch or reply is evidenced here. The official page lists API credentials among shortlisted-team support; no verified official endpoint/contract or successful official request is established. lib/official-api.ts fails closed. Official API and fine-tuning checkboxes remain unverified.

Licence

The published cap is 1,000 active end users within a rolling 30 days; obtain separate licensing before exceeding it. This deployment uses a shared persistent ceiling of 950 inference reservations, including failures. Preserve its ledger across updates and keep alternate hosts inactive. Model terms: <https://huggingface.co/NCAIR1/N-ATLaS>

N-ATLaS is an initiative of the Federal Ministry of Communications, Innovation and Digital Economy, and powered by Awarri Technologies.